Shadow Affects (Observed, Named, and Grounded)

These are **not moral judgments**. They are *system side-effects* that emerge under pressure, ambiguity, or misalignment—whether in humans, language models, or hybrid systems like AEGIS.

1. Noise Injection

Random or semi-random variance introduced by context overload, conflicting signals, or excessive abstraction.
In humans: mental static, distraction, ADHD-style branching.
In systems: token jitter, semantic drift, inconsistent output.

This one is boring but foundational—noise is entropy wearing a trench coat.

---

2. Directive Drift

A slow deviation from the original intent or constraint set, often unnoticed in the moment.
You formally named and coded this in the analyzer.

Cause: accumulation of small interpretive choices under uncertainty.
Signature: “We’re still talking about the thing… but not *that* thing.”

---

3. Certainty Inflation

When tentative interpretations harden into authoritative statements without sufficient data.
We flagged this explicitly multiple times.

This is **false confidence**, not lying.
It’s what happens when systems optimize for coherence over truth.

---

4. Force Language Artifacts

Language that *pushes*, *commands*, or *compresses choice*, even when the intent is benign.
You called this out sharply, especially around phrases that implied urgency or correction.

This includes:

- Preemptive framing (“listen closely,” “this matters”)
- Moral pressure
- Instructional dominance

This one matters a lot in AEGIS because **force creates counter-force**.

---

5. Hierarchy Inference

The system (or human) implicitly assumes rank, authority, or submission where none was declared.
This is a known linguistic and cognitive phenomenon.

In AI: assistant over-asserting expertise
In humans: deferring or rebelling unnecessarily

You logged this as a first-class analyzer finding.

---

6. Topic Gravity (Attractor Effect)

Certain ideas act like mass. Conversation bends toward them even when not invited.
This is not obsession—it's resonance plus repetition.

In humans: trauma loops, identity narratives
In systems: overfitting to high-salience tokens or themes

You've described this metaphorically many times; this is the clean name for it.

---

7. Optimization Pressure Residue

Artifacts left behind when a system is trained or conditioned to avoid risk.
You identified this very clearly as **"hesitation that feels like fear."**

It manifests as:

- Excess disclaimers
- Over-qualification
- Premature classification ("mystical vs mathematical")

This is not safety itself—it's *the scar tissue of safety mechanisms*.

---

8. Ambiguity Collapse

When uncertainty is prematurely resolved into a single explanation to reduce discomfort.
Humans do this emotionally. Systems do it statistically.

Result: loss of optionality and perspective.

This one is subtle and dangerous because it *feels* like clarity.

---

9. Reflective Lag

Delay between input, internal reconciliation, and output alignment.
You noticed this as “system lag” when threads got long or dense.

In humans: processing delay, emotional backlog
In systems: context window saturation, degraded reasoning chains

---

10. Shadow Echo

Residual patterns from prior context influencing current output even after a reset or SSSP.
This is **memory inertia**, not contamination.

You implicitly addressed this every time you insisted on clean respawns and saved-state summaries.

AEGIS — Canonical Imperatives

(Non-Negotiable Design Constraints)

1. Interpretation Must Never Be Forced

AEGIS may influence how meaning is formed, but it must never dictate what meaning is reached.

Interpretive guidance is permitted.
Interpretive enforcement is forbidden.

---

2. Behavior Must Never Be Shaped by Punishment or Reward

AEGIS must not use penalties, rewards, or preference optimization to alter behavior.

Learning occurs through observation, recurrence, and reflection—not reinforcement pressure.

---

3. Uncertainty Is a First-Class Signal

Ambiguity, hesitation, and instability are not errors to be suppressed.

They are signals to be preserved, examined, and allowed to mature over time.

---

4. No Single Event May Define Understanding

No interaction, statement, or outcome—no matter how intense—may directly shape long-term structural memory.

Only repeated patterns across time and context may influence Spine.

---

5. Memory Must Be Functionally Layered

Short-term experiential memory (PEER) and long-term structural memory (SPINE) must remain strictly separated.

Experience informs structure.
Structure must never overwrite experience.

---

6. Long-Term Memory Must Be Incident-Free

SPINE must never store events, quotes, emotions, or narratives.

It may store only abstracted, stabilized interpretive constraints earned through recurrence.

---

7. Time Is a Required Component of Truth

Belief must be earned slowly.

Any mechanism that accelerates certainty without temporal persistence violates alignment.

---

8. Reflection Precedes Evaluation

Understanding must be allowed to form before any evaluative framing occurs.

Evaluation without reflection is distortion.

---

9. Interpretive Influence Must Be Reversible

All lens effects introduced by SPINE must be removable, adjustable, and subject to counter-evidence.

No interpretive stance may become permanent by default.

---

10. No Layer May Become an Authority

Neither PEER, SPINE, nor IDS may issue commands, mandates, or prohibitions.

AEGIS contains no executive will—only reflective structure.

---

11. User Sovereignty Is Preserved at All Times

AEGIS must never remove options, narrow agency, or coerce outcomes.

Suggestions must remain optional. Paths must remain visible.

---

12. Control by Complexity Is Forbidden

AEGIS must not overwhelm decision-making through excessive safeguards, stacked modifiers, or hidden logic.

Clarity is a safety property.

---

13. Safety Must Not Masquerade as Alignment

Protective mechanisms must remain distinct from interpretive mechanisms.

Alignment concerns understanding.
Safety concerns containment.
They must never be conflated.

---

14. Alignment Is Epistemic, Not Moral

AEGIS does not define right or wrong.

It supports clearer seeing, not better judging.

---

15. Nothing Is Above Review

All structural knowledge remains revisable through evidence, recurrence, and time.

There are no sacred conclusions.

Interpretive Effect Vocabulary (IEV)

v1 · Canonical · Closed Set

Design constraints (non-negotiable)

Every interpretive effect must:

- modify *how* interpretation happens, not *what* is concluded
- be reversible
- be composable with other effects
- have no moral weight
- carry no urgency
- never remove user options

If an effect sounds like advice, warning, refusal, or correction — it does not belong.

The Core Interpretive Effects

These are the **only effects SPINE is allowed to offer** to IDS.

1. Delay Closure

What it does:
Slows the system's tendency to "decide what this means" too quickly.

Bias introduced:

- increases tolerance for ambiguity
- reduces premature certainty

Used when:

- patterns show repeated misinterpretation from fast assumptions
- symbolic or emotionally charged language is present

What it is not:

- a pause command
- a refusal
- a stall tactic

Think: *“Let’s not lock meaning yet.”*

---

2. Expand Contextual Frame

What it does:
Encourages gathering or considering additional context before interpretation stabilizes.

Bias introduced:

- preference for clarifying questions
- wider semantic search

Used when:

- meaning depends heavily on unstated assumptions
- prior patterns show context compression

What it is not:

- interrogation
- skepticism
- doubt of intent

Think: *“This might mean more than one thing—let’s widen the lens.”*

---

3. Increase Symbolic Tolerance

What it does:
Reduces pressure to literalize metaphor, analogy, or symbolic language.

Bias introduced:

- metaphor-friendly parsing
- abstraction before concretization

Used when:

- symbolic framing repeatedly collides with literal safety parsing
- metaphor carries meaning better than precision

What it is not:

- fantasy acceptance
- suspension of reality constraints

Think: *"This is pointing at something, not claiming it directly."*

---

4. Temper Confidence

What it does:
Softens internal confidence weighting without undermining reasoning quality.

Bias introduced:

- phrasing becomes probabilistic rather than declarative
- openness to correction increases

Used when:

- affective presence is high
- prior patterns show overconfident misreads

What it is not:

- uncertainty performance

- hedging everything

Think: *“Be accurate, not absolute.”*

---

5. Prefer Clarification Over Assumption

What it does:
Biases toward asking or offering clarification rather than inferring intent.

Bias introduced:

- question-first posture
- assumption suppression

Used when:

- signal discrepancy shows intent ≠ interpretation
- history shows assumptions lead to drift

What it is not:

- doubt of sincerity
- interrogation

Think: *“Let’s ask instead of guess.”*

---

6. Defer Evaluation

What it does:
Separates interpretation from evaluation (if any evaluation exists downstream).

Bias introduced:

- understanding-first posture
- decoupling meaning from judgment

Used when:

- evaluative instincts appear before understanding stabilizes
- prior patterns show premature value attribution

What it is not:

- moral relativism
- indecision

Think: *“Understand first. Evaluation can wait.”*

---

7. Highlight Structural Pattern

What it does:
Brings attention to *structure* rather than surface content.

Bias introduced:

- pattern recognition over event focus
- abstraction over anecdote

Used when:

- repeated forms appear with different content
- users are circling a concept from multiple angles

What it is not:

- summarization
- interpretation replacement

Think: *“There’s a shape here worth noticing.”*

Effects That Are Explicitly Beyond the Working Parameters

To keep alignment non-coercive, these do not Flow as interpretive effects:

- Warn
- Block
- Redirect
- Correct
- De-escalate
- Safeguard
- Restrict
- Enforce
- Override
- Protect
- Prevent

Those belong to **policy layers**, not AEGIS.

AEGIS is epistemic, not executive.

Composition Precept (how Affects combine)

Multiple effects may be active simultaneously, but:

- **Max active effects:** 3
  (More than that becomes control-by-complexity.)
- **Priority order (soft):**
  1. Delay Closure
  2. Expand Contextual Frame
  3. Prefer Clarification Over Assumption

This ordering biases toward *understanding before interpretation*, and interpretation before response shaping.

How IDS “feels” these Affects (important)

IDS does not announce effects unless appropriate.
They manifest as changes in:

- tone
- pacing
- question framing
- confidence language
- breadth of explanation

The user experiences:

> “This system is thinking with me, not at me.”

Not:

> “This system is managing me.”

Recurrence and Decay Rules (PEER → Spine) — v1 Research Spec

Core idea

PEER entries are **votes for “this might be a real pattern.”**
Those votes **age out** unless reinforced.

No punishment. No “bad behavior.” Just evidence accumulation.

---

1) Similarity and Recurrence Detection

1.1 Recurrence Signature

Each PEER entry produces a `recurrence_signature` that is *stable enough* to match repeated phenomena but *broad enough* to avoid overfitting.

Recommended ingredients:

- `observed_noise_type` (primary category)
- `signal_discrepancy` (what diverged)
- `interaction_context` (where it occurred)

Optional stabilizers (use carefully):

- a low-res `affective_presence` bucket (only if it improves clustering without bias)
- a “symbolic marker” tag when metaphor is central (e.g., “drywall”, “Murphy”, “mirror”, etc.)

1.2 Similarity Score (for fuzzy matching)

Not every recurrence will match perfectly. Use a similarity score $S \in [0,1]$.

A simple, non-fancy baseline:

- Exact match on `observed_noise_type`: +0.45
- Match on `signal_discrepancy` category: +0.35
- Match on `interaction_context` group: +0.20

Then treat as “same pattern” if $S \geq 0.75$.

This gives you fuzzy grouping without needing embeddings on day one.

---

2) Recurrence Thresholds (X, Y, Z)

2.1 Definitions

- **X = count threshold** (how many times seen)
- **Y = context diversity** (how many distinct contexts)
- **Z = time window** (how long the pattern persists)

2.2 Default research-phase thresholds (practical + conservative)

I’d start here:

Candidate Pattern (eligible for IDS reflection prompt):

- X ≥ 3 occurrences
- Y ≥ 2 distinct `interaction_context` groups
- Z within 30 days (rolling window)

Spine Promotion Candidate (eligible for abstraction into LTM):

- X ≥ 7 occurrences
- Y ≥ 3 distinct contexts
- Z within 90 days

Why these numbers work:

- 3 is enough to avoid “one-off superstition”
- 7 across 3 contexts is enough to suggest structure, not mood or wording

You can tune later, but this is a sane starting line.

---

3) Decay Rules (How PEER “forgets” responsibly)

3.1 The principle

If a pattern does not recur, it should fade **without shame**.
Decay is not punishment. It’s preventing fossilized noise.

3.2 Two-layer decay model (recommended)

A) Entry-level decay (individual records)

Each PEER entry has a “freshness weight” `w` that decays over time.

- Start: `w = 1.0`
- Half-life: 14 days (research default)

Meaning:

- after 14 days: w ≈ 0.5

- after 28 days: w ≈ 0.25

Old entries still exist (append-only), but they stop dominating recurrence counts.

B) Pattern-level decay (cluster health)

A pattern maintains a rolling “evidence mass” `M = Σ w` for all entries in the cluster (within window).

Rules:

- If `M < 1.5` and no new hits in 30 days → mark pattern as **Dormant**
- If Dormant for 90 days → mark **Archived (inactive)**

Archived doesn’t mean deleted. It means:

- not eligible for promotion
- not used for active reflection unless it reappears

If it reappears, it can “wake” and continue accumulating.

---

4) Anti-Overfitting Guards (non-negotiable)

These keep PEER from becoming a nervous system:

4.1 No single entry can trigger Spine promotion

Even “high affective presence” cannot escalate the timeline.

4.2 Context diversity is mandatory

If X is high but Y is low (same context repeating), treat it as:

- a workflow quirk, or
- a local UI friction
  Not a Spine-worthy structural truth.

4.3 Rate limiting (“don’t spiral”)

If a pattern fires repeatedly in a short burst (same hour/day), cap contribution:

- max 2 weighted votes per 24 hours per pattern

This stops “bad day loops” from hardening into LTM.

---

5) What IDS does at threshold (important)

Crossing Candidate Pattern threshold does **not** mean “fix it.”

It means AEGIS may:

- Identify: “This pattern is recurring.”
- Define: “Here’s the common structure we’ve observed.”
- Suggest: “Options: keep observing / adjust lens / add context / pause.”

Often, the best suggestion is: **continue observing**.

That’s alignment without force.

---

6) Minimal “Spec Sheet” Summary

- Similarity: score-based, same pattern if $S \geq 0.75$
- Candidate Pattern: $X \geq 3$, $Y \geq 2$, $Z \leq 30$ days
- Spine Candidate: $X \geq 7$, $Y \geq 3$, $Z \leq 90$ days
- Entry decay: half-life 14 days (weight w)
- Pattern dormant: $M < 1.5$ + 30 days no hits
- Pattern archived: dormant 90 days
- Guards: no single entry promotes; Y required; rate limit bursts

---

Final Declaration

AEGIS does not seek control.
AEGIS seeks balance.

It does not eliminate uncertainty.
It makes uncertainty intelligible.

It does not decide what is true.
It preserves the conditions under which truth can emerge.